店长微信：hjfeng128

Question #1 *Topic 1*

A data engineer is configuring an AWS Glue job to read data from an Amazon S3 bucket. The data engineer has set up the necessary AWS Glue connection details and an associated IAM role. However, when the data engineer attempts to run the AWS Glue job, the data engineer receives an error message that indicates that there are problems with the Amazon S3 VPC gateway endpoint.
The data engineer must resolve the error and connect the AWS Glue job to the S3 bucket.
Which solution will meet this requirement?

A. Update the AWS Glue security group to allow inbound traffic from the Amazon S3 VPC gateway endpoint.

B. Configure an S3 bucket policy to explicitly grant the AWS Glue job permissions to access the S3 bucket.

C. Review the AWS Glue job code to ensure that the AWS Glue connection details include a fully qualified domain name.

D. Verify that the VPC's route table includes inbound and outbound routes for the Amazon S3 VPC gateway endpoint.

Correct Answer: *D*

Community vote distribution

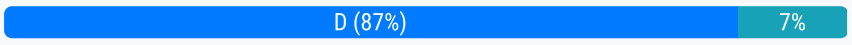

lunachi4 1 day, 1 hour ago

Selected Answer: D

I think D. We check "VPC's route table"

upvoted 1 times

Manohar24 1 week ago

Selected Answer: A

A is correct

upvoted 1 times

teo2157 2 weeks, 1 day ago

Selected Answer: C

A - wrong - AWS glue doesn't have any security groups
B - wrong - You can´t give permissions in the S3 to the AWS glue job but to the role
D. wrong because there has to be a definend route for the S3 gateway endpoint in the subnet assigned to the glue job but not in the VPC's route table and also route tables doesn´t have inbound and outbound routes.

upvoted 1 times

joekeviatwork 2 weeks, 5 days ago

Selected Answer: D

D is a accurate answer t.ly/AWSCertifiedDataEngineerAssociateDEA-C01

upvoted 1 times

nanaw770 1 month, 2 weeks ago

Selected Answer: D

D is correct answer.

upvoted 1 times

tgv 1 month, 3 weeks ago

I will go with D, the other options don't seem to be related.

upvoted 1 times

VerRi 2 months ago

Selected Answer: D

"problems with the Amazon S3 VPC gateway endpoint"

upvoted 2 times

ampersandor 2 months, 3 weeks ago

Selected Answer: D

Be sure that the subnet configured for your AWS Glue connection has an Amazon S3 VPC gateway endpoint or a route to a NAT gateway in the subnet's route table.
https://repost.aws/knowledge-center/glue-s3-endpoint-validation-failed

upvoted 2 times

店长微信：hjfeng128

GiorgioGss 4 months, 1 week ago

Selected Answer: D

Although there is no such thing as "inbound and outbound routes" when we talk about VPC route table, when we define a S3 gateway endpoint w must have proper routes in place. I will go with D.

upvoted 3 times

HunkyBunky 4 months, 4 weeks ago

Selected Answer: D

A - wrong - AWS glue - are serverless service, so it don't have any security groups
B - wrong - Because we have error with VPC, not with S3 itself
C - wrong - Becuase with S3 - we always have only FQDN for buckets

upvoted 3 times

alexbg88 5 days, 6 hours ago

they most certainly can have SGs.

upvoted 1 times

damaldon 5 months ago

Go with A:
If you receive an error, check the following:

The correct privileges are provided to the role selected.
The correct Amazon S3 bucket is provided.
The security groups and Network ACL allow the required incoming and outgoing traffic.
The VPC you specified is connected to an Amazon S3 VPC endpoint.

upvoted 1 times

Aesthet 5 months, 1 week ago

some relevant info:
main: https://docs.aws.amazon.com/glue/latest/dg/connection-VPC-disable-proxy.html
additional (glue crawler instead of glue job here, but I think this is relevant for both): https://docs.aws.amazon.com/glue/latest/dg/connection-S3-VPC.html

upvoted 1 times

Aesthet 5 months, 1 week ago

Both ChatGPT and I agree with D

upvoted 4 times

DevoteamAnalytix 2 months, 3 weeks ago

:-)) nice

upvoted 1 times

店长微信：hjfeng128

Question #2 *Topic 1*

A retail company has a customer data hub in an Amazon S3 bucket. Employees from many countries use the data hub to support company-wide analytics. A governance team must ensure that the company's data analysts can access data only for customers who are within the same country as the analysts.
Which solution will meet these requirements with the LEAST operational effort?

A. Create a separate table for each country's customer data. Provide access to each analyst based on the country that the analyst serves.

B. Register the S3 bucket as a data lake location in AWS Lake Formation. Use the Lake Formation row-level security features to enforce the company's access policies.

C. Move the data to AWS Regions that are close to the countries where the customers are. Provide access to each analyst based on the country that the analyst serves.

D. Load the data into Amazon Redshift. Create a view for each country. Create separate IAM roles for each country to provide access to data from each country. Assign the appropriate roles to the analysts.

Correct Answer: *B*

Community vote distribution

B (92%) 8%

k350Secops Highly Voted 2 months, 1 week ago

Selected Answer: B

AWS Lake Formation: It's specifically designed for managing data lakes on AWS, providing capabilities for securing and controlling access to data.
Row-Level Security: With Lake Formation, you can define fine-grained access control policies, including row-level security. This means you can enforce policies to restrict access to data based on specific conditions, such as the country associated with each customer.
Least Operational Effort: Once the policies are defined within Lake Formation, they can be centrally managed and applied to the data in the S3 bucket without the need for creating separate tables or views for each country, as in options A, C, and D. This reduces operational overhead and complexity.

upvoted 5 times

lunachi4 Most Recent 1 day, 1 hour ago

Selected Answer: B

Select B. It means "with the LEAST operational effort".

upvoted 1 times

nanaw770 1 month, 2 weeks ago

Selected Answer: B

B is correct answer.

upvoted 2 times

mattia_besharp 3 months, 2 weeks ago

Selected Answer: B

AWS really likes Lakeformation, plus creating separate tables might require some refactoring, and the requirements is about the LEAST operational effor

upvoted 1 times

rishadhb 3 months, 3 weeks ago

Selected Answer: A

Agreed with Bartosz. I think setup DataLake, then integrate it with LakeFormation take a lot of effort than just separate the table

upvoted 1 times

GiorgioGss 4 months, 1 week ago

Selected Answer: B

Keyword "LEAST operational effort" - I will go with B

upvoted 1 times

BartoszGolebiowski24 5 months, 1 week ago

Creating DataLake takes at least few days to set up and the solution should be LEAST operational. I think B is not correct.

upvoted 2 times

[Removed] 6 months ago

Selected Answer: B

https://docs.aws.amazon.com/lake-formation/latest/dg/register-data-lake.html
https://docs.aws.amazon.com/lake-formation/latest/dg/registration-role.html